



A Next Generation Monitoring System



Who am I?



Richard Steinbrück
MC Store

DevOps, Automation, Puppet, Ansible, Monitoring ...

Why monitor?

There are many reasons to monitor a system, including:

Analyzing long-term trends

- How big is my database and how fast is it growing?

Comparing over time or experiment groups

- How much better is my memcache hit rate with an extra node? Is my site slower than it was last week?

Alerting

- Something is broken, and somebody needs to fix it right now!

Building dashboards

- Dashboards should answer basic questions about your service

Conducting ad hoc retrospective analysis

- Our latency just shot up; what else happened around the same time?

How Monitoring looks today?

Nagios/Icinga/CheckMK/Zabbix

Configuration on Master
Monitoring-Agents with Script Extensions
Integrated Dashboards



Prometheus

Inspired by Google's Borgmon monitoring system


Started in 2012 by ex-Googleers working in Soundcloud as an open source project.

Mainly written in Go. Version 1.0 released in 2016



500+ Companies using it including Digital Ocean, CoreOS, Airtame and Jodel



Selling Points

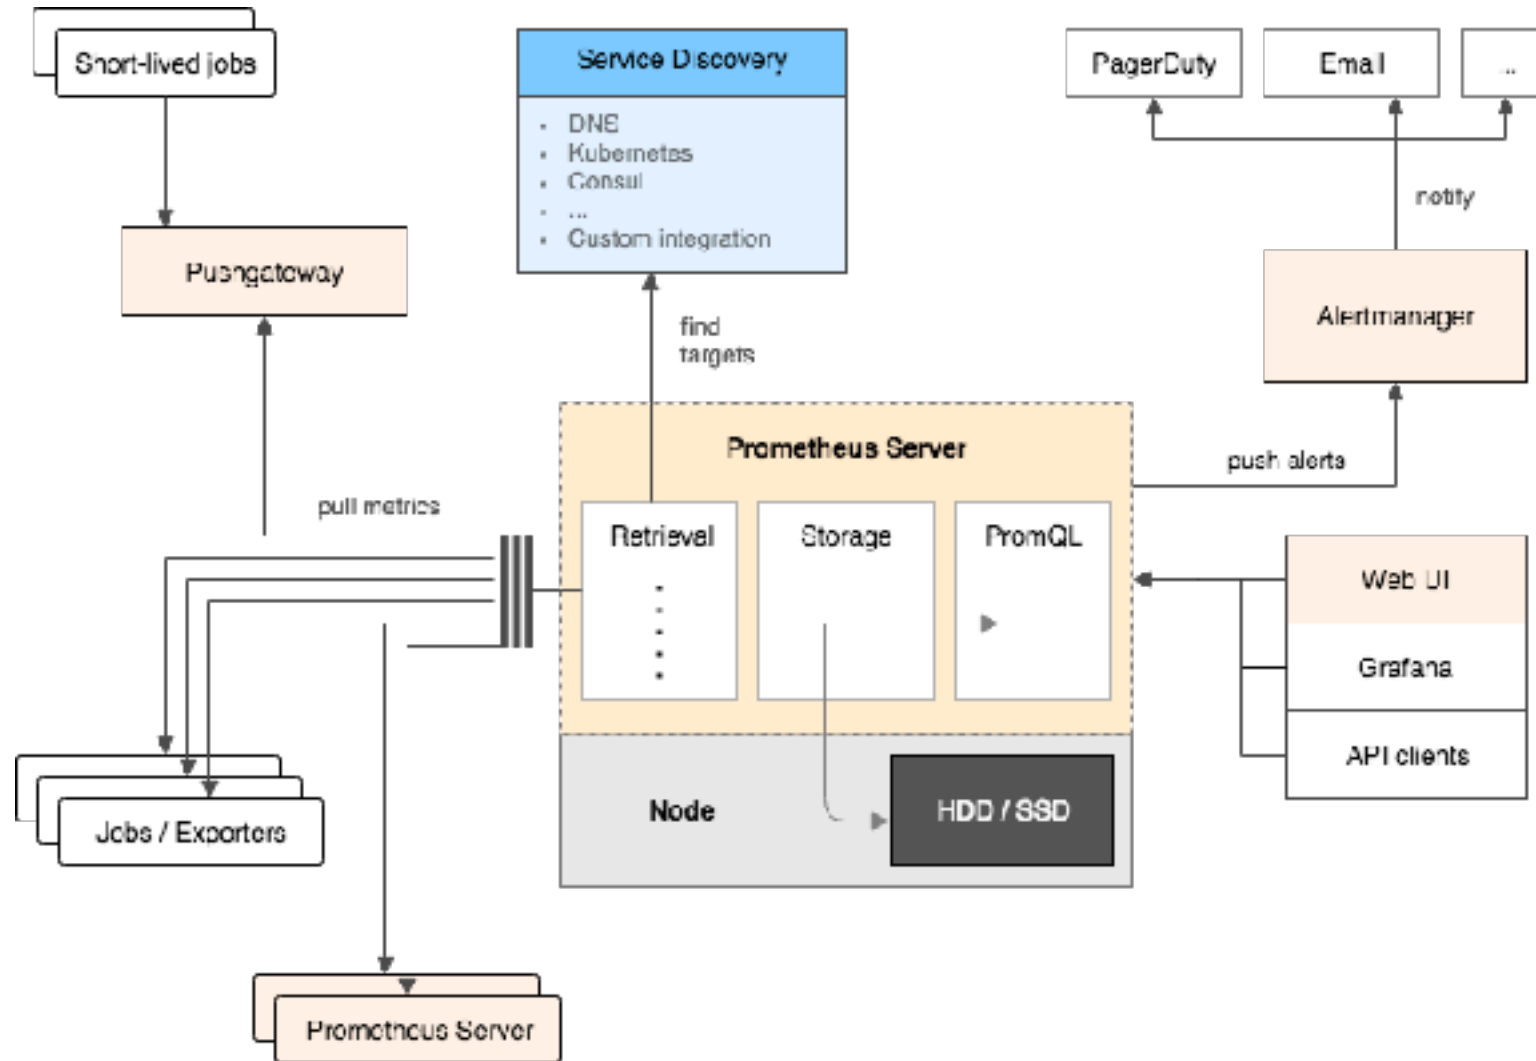


Blackbox vs. Whitebox Monitoring



DO
ONE THING
AND DO IT WELL!

Architecture - Overview



Architecture - Exporter

node_cpu{instance="ing04gmispup01d.media-saturn.com:9100"}

Execute

- insert metric at cursor -

Graph

Console

Element	Value
node_cpu{cpu="cpu0",group="linux",instance="ing04gmispup01d.media-saturn.com:9100",job="GDP-PROD",mode="guest"}	0
node_cpu{cpu="cpu0",group="linux",instance="ing04gmispup01d.media-saturn.com:9100",job="GDP-PROD",mode="guest_nice"}	0
node_cpu{cpu="cpu0",group="linux",instance="ing04gmispup01d.media-saturn.com:9100",job="GDP-PROD",mode="idle"}	4198517.54
node_cpu{cpu="cpu0",group="linux",instance="ing04gmispup01d.media-saturn.com:9100",job="GDP-PROD",mode="iowait"}	14273.47

← → ↺

ing04gmispup01d.media-saturn.com:9100/metrics

```
# TYPE node_boot_time gauge
node_boot_time 1.516711442e+09
# HELP node_context_switches Total number of context switches.
# TYPE node_context_switches counter
node_context_switches 1.4470104385e+10
# HELP node_cpu Seconds the cpus spent in each mode.
# TYPE node_cpu counter
node_cpu{cpu="cpu0",mode="guest"} 0
node_cpu{cpu="cpu0",mode="guest_nice"} 0
node_cpu{cpu="cpu0",mode="idle"} 4.19841085e+06
node_cpu{cpu="cpu0",mode="iowait"} 14273.37
```

Easy to Integrate

For DEV's: Go, Java or Scala, Python, Ruby, Bash, C++, Common Lisp, Elixir, Erlang, Haskell, .NET / C#, Node.js, PHP, Rust



For OP's: Node_Exporter, WMI_Exporter, Nginx, Apache, MySQL, Postgres, AWS, Google Stackdrive Varnish, RabbixMQ, IPMI, MongoDB, JMX, SNMP, Jenkins, Docker, Nagios ... and many more ;)



Architecture - Services Discovery

Microsoft Azure

Services Discovery configurations allow retrieving scrape targets from ...
azure, consul, dns, ec2, openstack, file, gce, kubernetes, marathon, nerve, serverset, triton ... and many more



kubernetes
by Google



← → ↻ ⓘ ing04gmispup01d.media-saturn.com:9898/gmsrest/prometheussd/targets

```
- targets:
- TQA-WIN-D00938.media-saturn.com:9182
- TQA-WIN-D00940.media-saturn.com:9182
- TQA-WIN-D00941.media-saturn.com:9182
- TQA-WIN-D01026.media-saturn.com:9182
- TQA-WIN-D01027.media-saturn.com:9182
- TQA-WIN-D01046.media-saturn.com:9182
- TQA-WIN-D01117.media-saturn.com:9182
- TQA-WIN-D01128.media-saturn.com:9182
- TQA-WIN-D01133.media-saturn.com:9182
- TQA-WIN-D01141.media-saturn.com:9182
- TQA-WIN-D01142.media-saturn.com:9182
- TQA-WIN-D01154.media-saturn.com:9182
```

Query your data

Top 5 Docker images by CPU

```
topk(5,  
  sum by (image)(  
    rate(container_cpu_usage_seconds_total{  
      image=~"prod.docker-hub.media-saturn.com/storeops/.*"}[5m])  
  )  
)
```

Disk Will Fill in 4 Hours


```
predict_linear(node_filesystem_free[1h], 4*3600) < 0
```

Which of your Alerts have been firing the most?

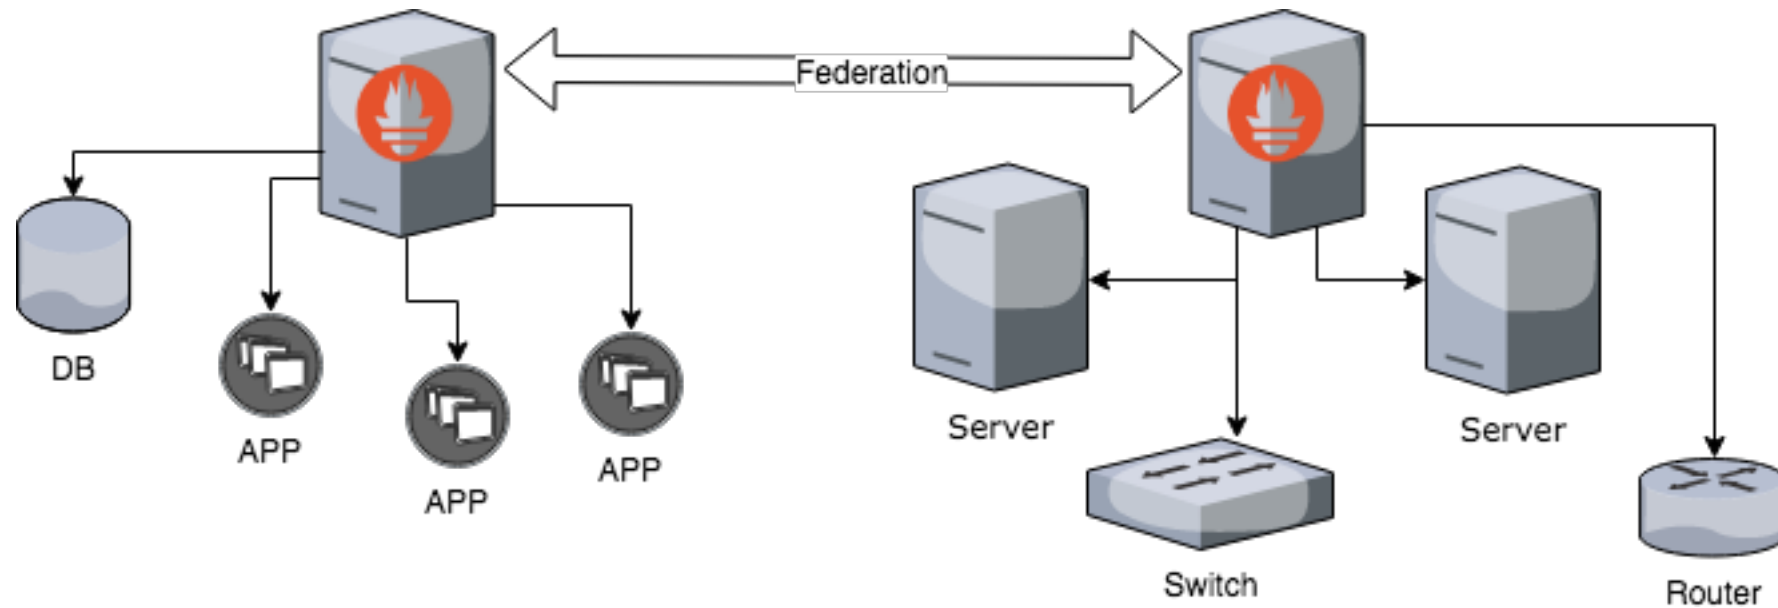
```
sum(  
  sort_desc(sum_over_time(ALERTS{alertstate=firing}[24h]))  
  ) by (alertname)
```

Architecture - Pushgateway

```
echo "some_metric 3.14" | curl --data-binary @- http://pushgateway:9091/metrics/job/some_job
```



Architecture - Federation



Federation allows a Prometheus server to scrape selected time series from another Prometheus server.

Alerting

- alert: InstanceDown
 expr: up == 0
 for: 1m
 labels:
 severity: critical
 annotations:
 description: "Instance {{ \$labels.instance }} is down"
- alert: WebAppIsDown
 expr: probe_http_status_code != 200
 for: 30s
 labels:
 severity: warning
 annotations:
 description: "Webpage {{ \$labels.instance }} is Down!"



QUESTIONS?

Resources

Slides und Code: https://github.com/steinbrueckri/MS_TechSummit_2018

Query Examples: <https://github.com/infinityworks/prometheus-example-queries>

<https://prometheus.io/>